



A Deeper Dive

Soumith Chintala
Facebook AI Research


Overview

- Tensors and Storages
- Neural Networks




Tensors and Storages

- Tensor = n-dimensional array
- Row-major in memory



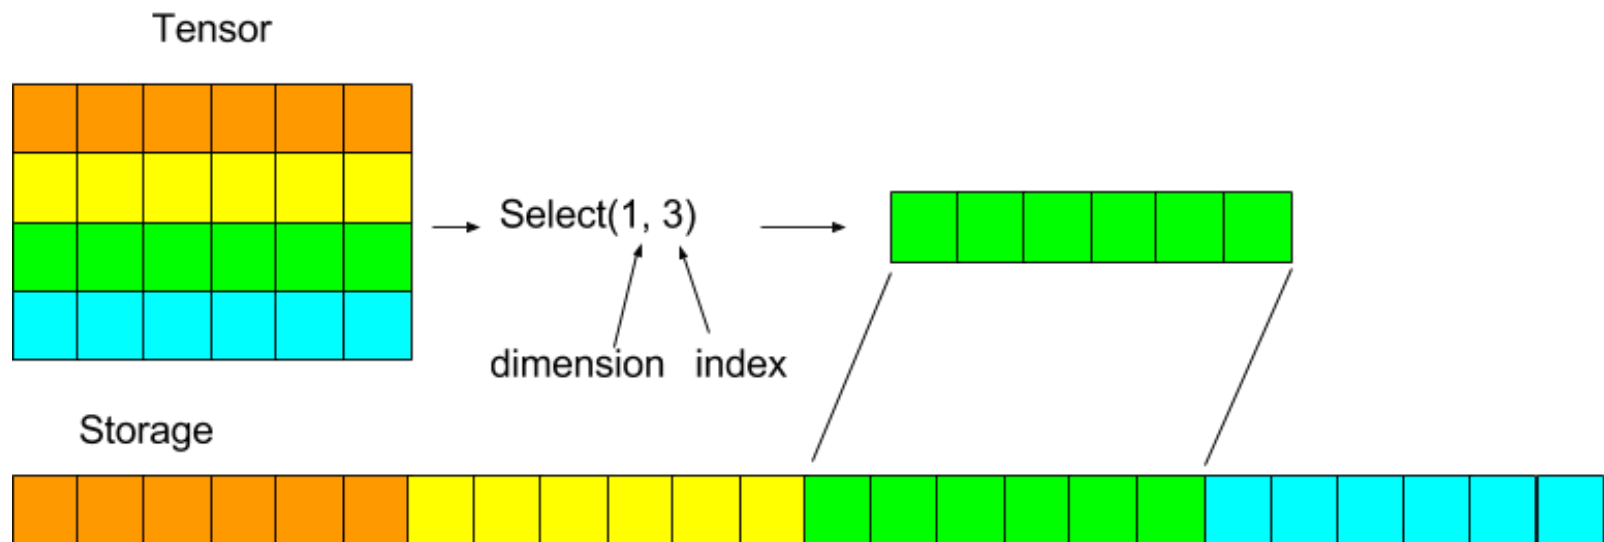
Tensors and Storages

- Tensor = n-dimensional array
- Row-major in memory



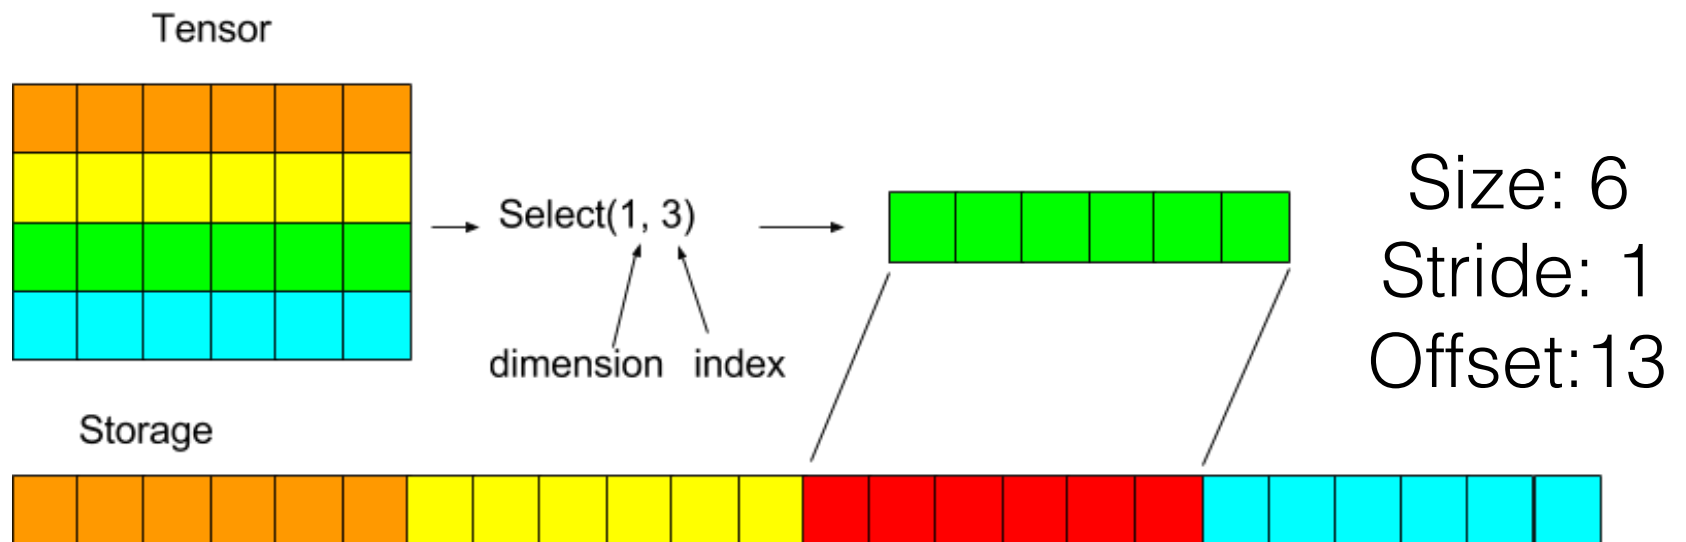
Tensors and Storages

- Tensor = n-dimensional array
- 1-indexed



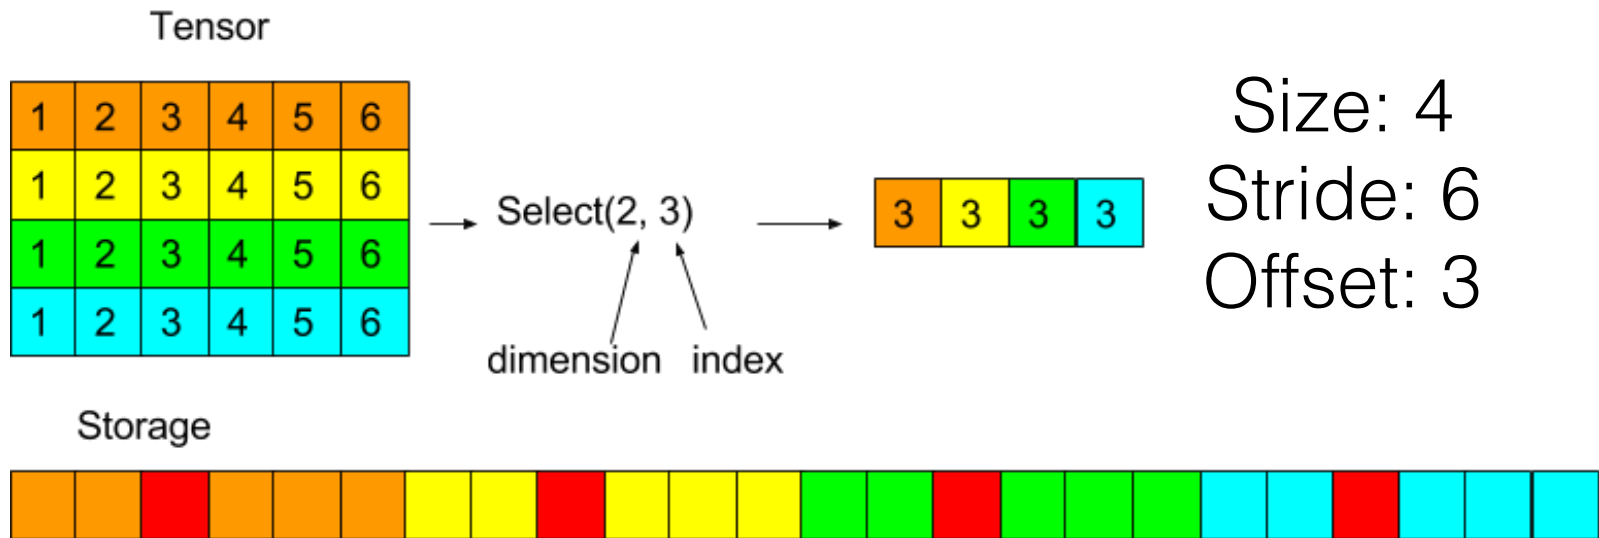
Tensors and Storages

- Tensor = n-dimensional array
- Tensor: size, stride, storage, storageOffset



Tensors and Storages

- Tensor = n-dimensional array
- Tensor: size, stride, storage, storageOffset




Tensors and Storages

```
In [1]: require 'torch';
```

```
In [2]: a = torch.DoubleTensor(4, 6) -- DoubleTensor, uninitialized memory
        a:uniform() -- fills a with uniform noise with mean = 0, stdv = 1
```

```
In [3]: print(a)
```

```
Out[3]: 0.4332 0.5716 0.5750 0.8167 0.1997 0.6187
         0.7775 0.3575 0.0749 0.4028 0.0532 0.4481
         0.5088 0.1795 0.6948 0.5700 0.7679 0.6176
         0.9225 0.7270 0.2223 0.1087 0.2717 0.8853
         [torch.DoubleTensor of size 4x6]
```



Tensors and Storages

```
In [1]: require 'torch';
```

```
In [2]: a = torch.DoubleTensor(4, 6) -- DoubleTensor, uninitialized memory  
a:uniform() -- fills a with uniform noise with mean = 0, stdv = 1
```

```
In [3]: print(a)
```

```
Out[3]:  0.4332  0.5716  0.5750  0.8167  0.1997  0.6187  
         0.7775  0.3575  0.0749  0.4028  0.0532  0.4481  
         0.5088  0.1795  0.6948  0.5700  0.7679  0.6176  
         0.9225  0.7270  0.2223  0.1087  0.2717  0.8853  
[torch.DoubleTensor of size 4x6]
```

```
In [4]: b = a:select(1, 3)
```

```
In [5]: print(b)
```

```
Out[5]:  0.5088  
         0.1795  
         0.6948  
         0.5700  
         0.7679  
         0.6176  
[torch.DoubleTensor of size 6]
```

Tensors and Storages

Underlying storage is shared

```
In [6]: b:fill(3);
```

```
In [7]: print(b)
```

```
Out[7]:  3
         3
         3
         3
         3
         3
         3
[torch.DoubleTensor of size 6]
```

```
In [8]: print(a)
```

```
Out[8]:  0.4332  0.5716  0.5750  0.8167  0.1997  0.6187
         0.7775  0.3575  0.0749  0.4028  0.0532  0.4481
         3.0000  3.0000  3.0000  3.0000  3.0000  3.0000
         0.9225  0.7270  0.2223  0.1087  0.2717  0.8853
[torch.DoubleTensor of size 4x6]
```



Tensors and Storages

- 150+ Tensor functions
 - Linear algebra
 - Convolutions
 - BLAS
 - Tensor manipulation
 - Narrow, index, mask, etc.
 - Logical operators
- Fully documented: <https://github.com/torch/torch7/tree/master/doc>
- Inline help!



Tensors and Storages

- Inline help

```
In [10]: ?torch.cmul
```

```
Out[10]: ++++++
[res] torch.cmul([res,] tensor1, tensor2)
```

Element-wise multiplication of `tensor1` by `tensor2` .
The number of elements must match, but sizes do not matter.

```
> x = torch.Tensor(2, 2):fill(2)
> y = torch.Tensor(4):fill(3)
> x:cmul(y)
> = x
 6  6
 6  6
[torch.DoubleTensor of size 2x2]
```

`z = torch.cmul(x, y)` returns a new Tensor .

`torch.cmul(z, x, y)` puts the result in `z` .

`y:cmul(x)` multiplies all elements of `y` with corresponding elements of `x` .

`z:cmul(x, y)` puts the result in `z` .

```
+++++
```



Tensors and Storages

- GPU support for all operations:
 - require 'cutorch'
 - torch.CudaTensor = torch.FloatTensor on GPU
- Fully multi-GPU compatible

```
In [ ]: require 'cutorch'  
a = torch.CudaTensor(4, 6):uniform()  
b = a:select(1, 3)  
b:fill(3)
```




Training Neural Networks



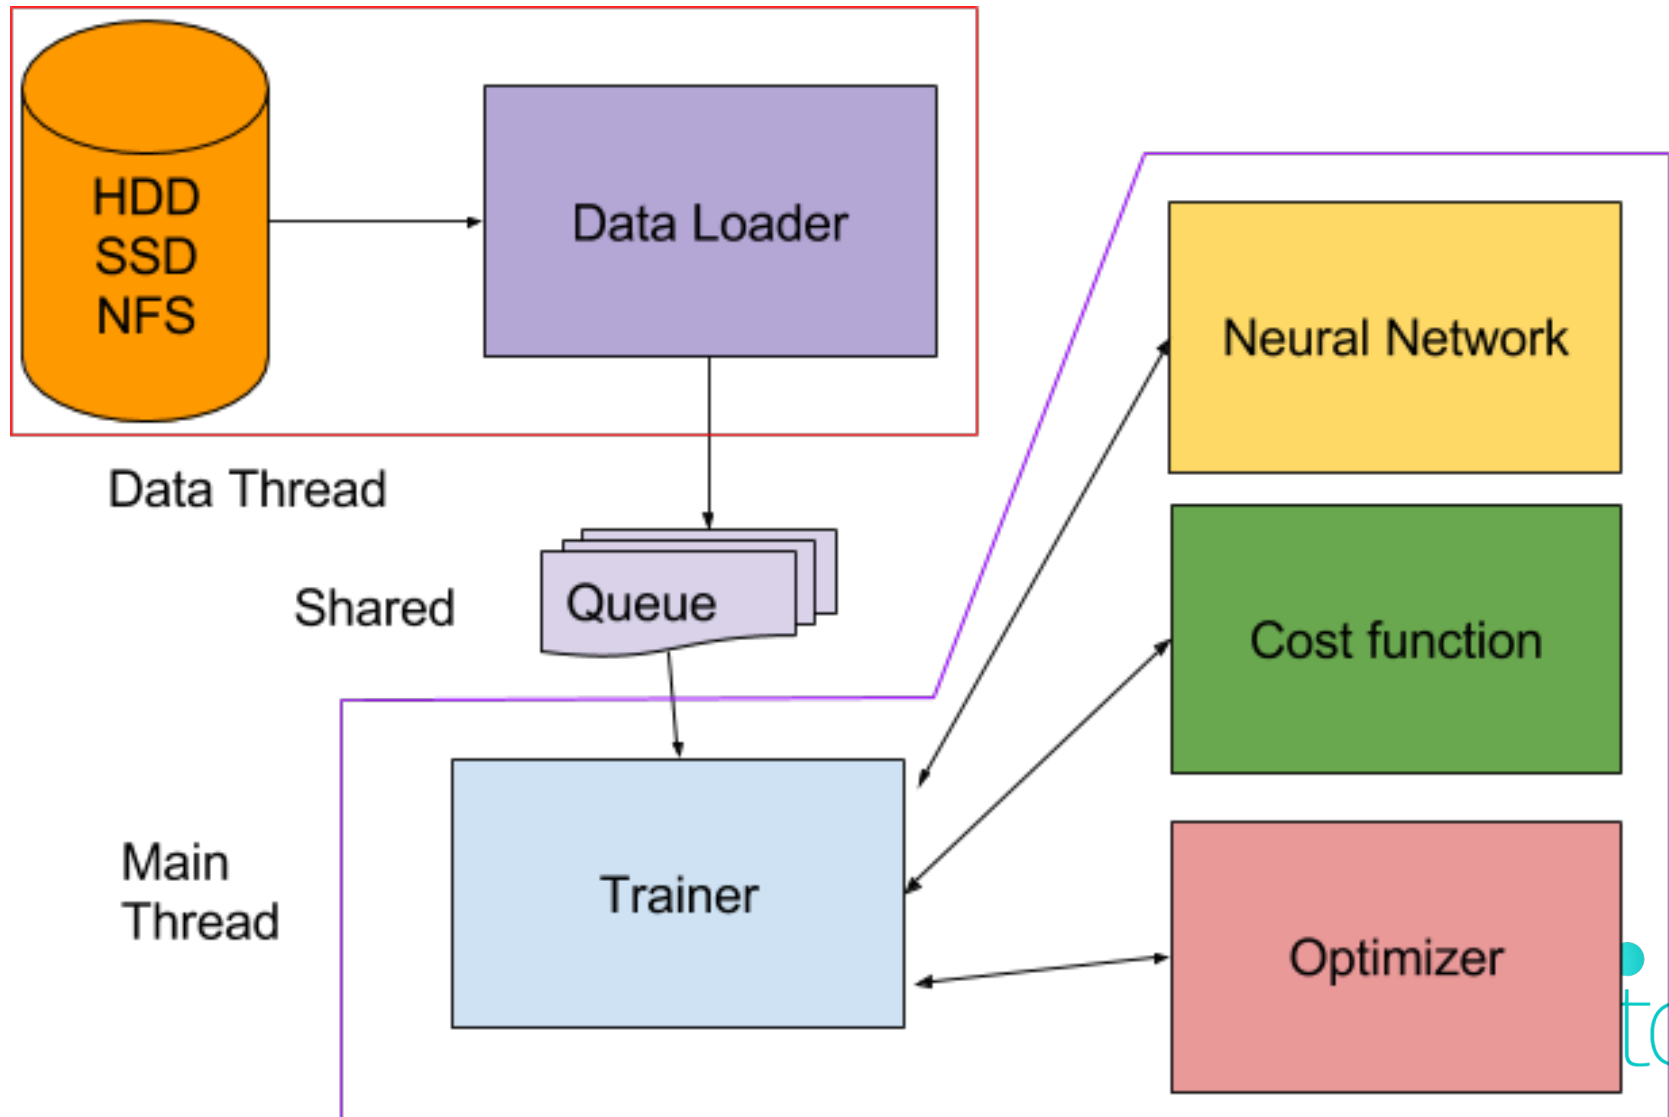
Training cycle

Moving parts



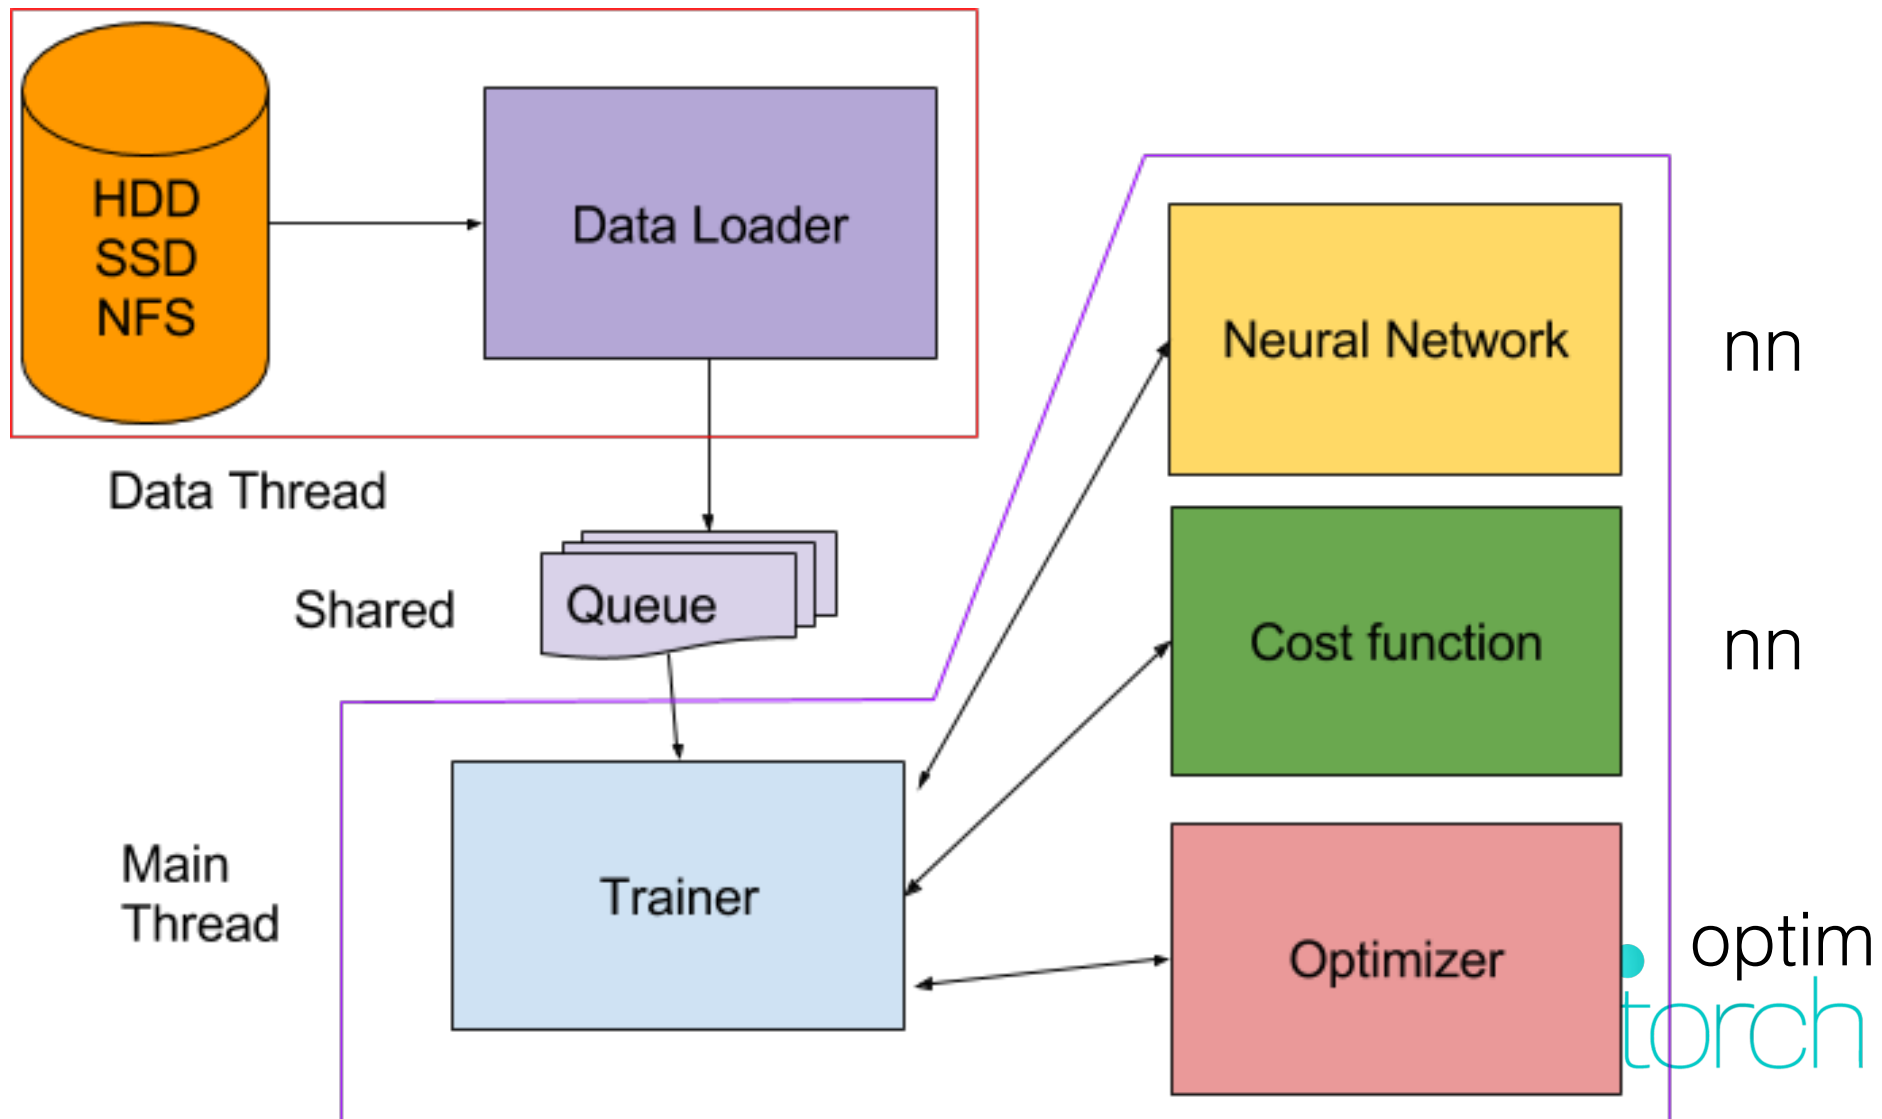
Training cycle

Moving parts

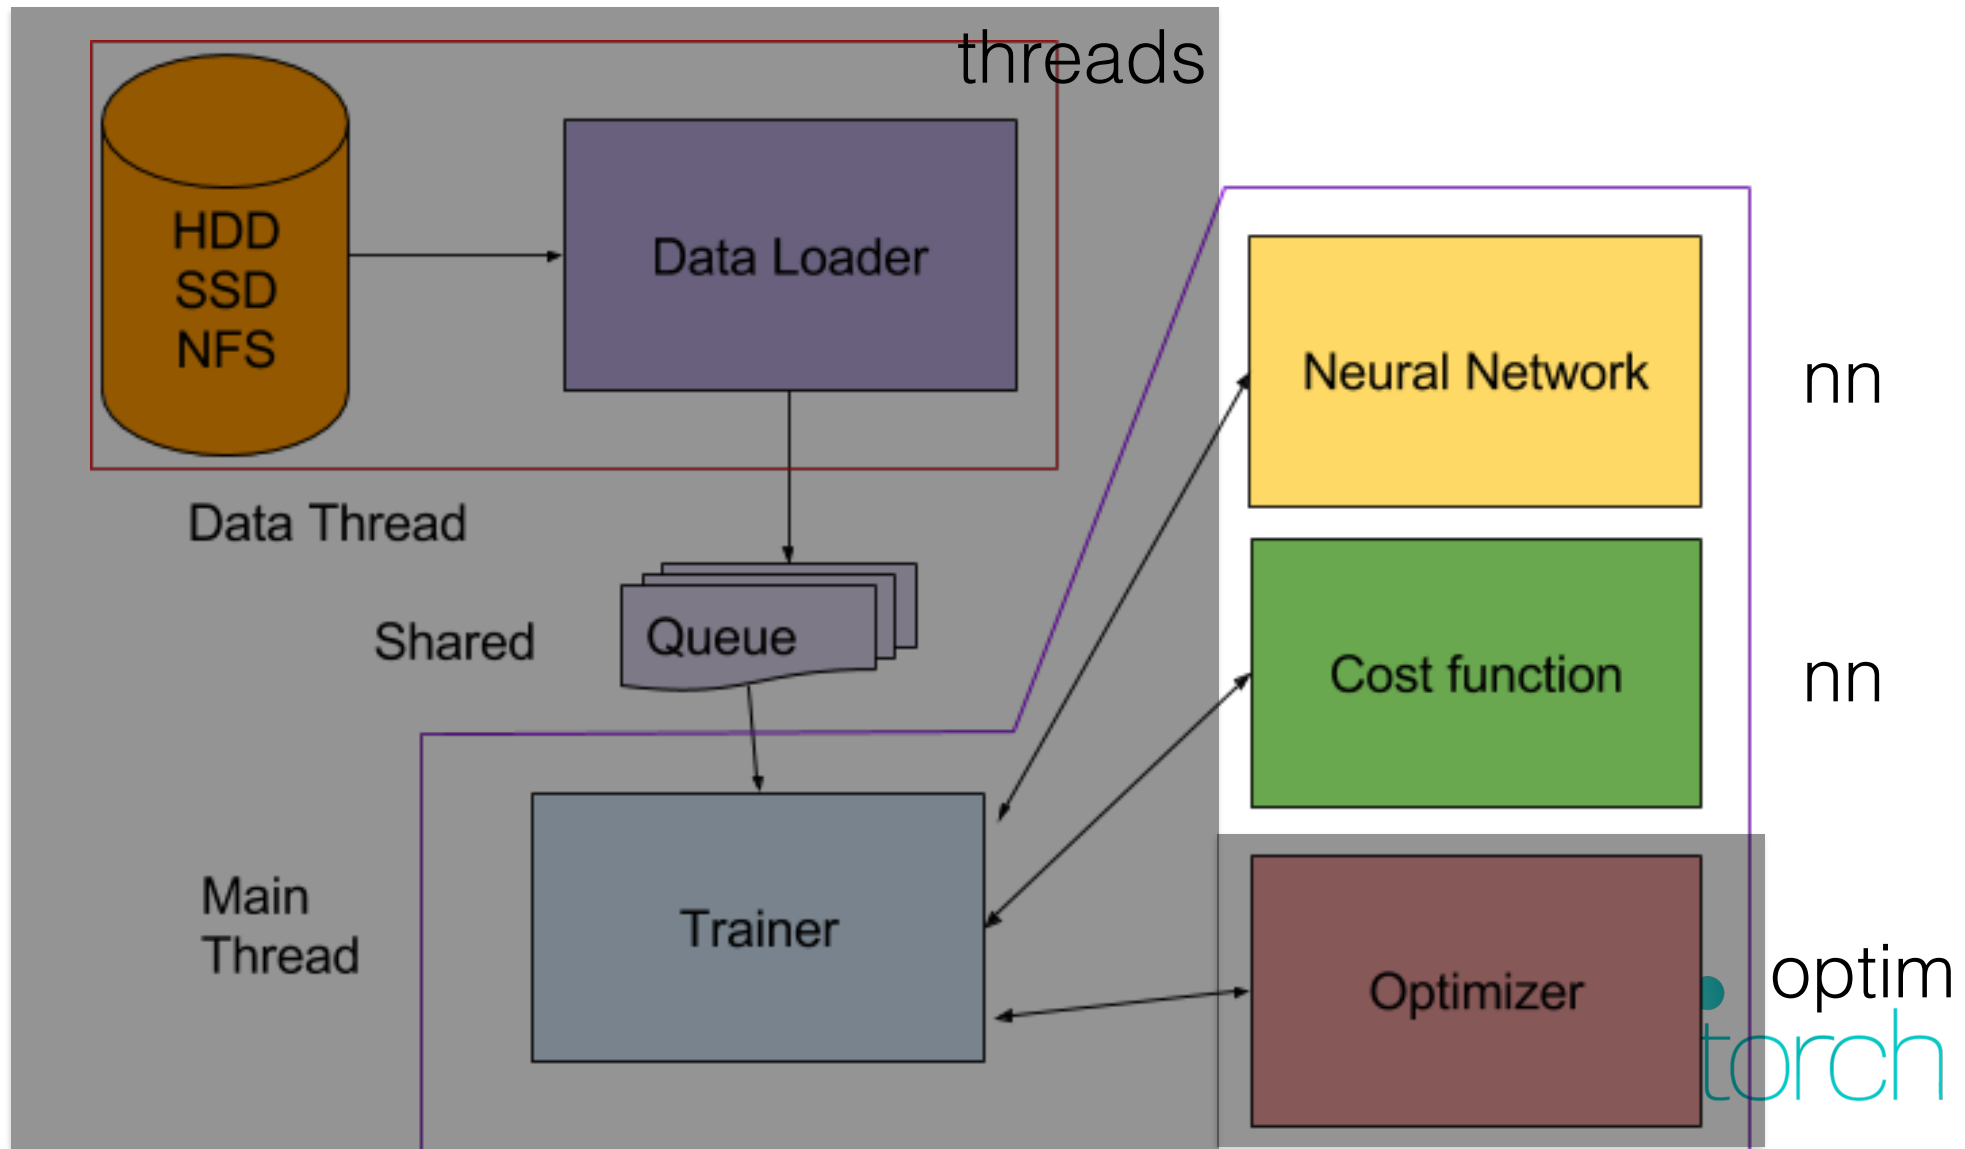


Training cycle

threads



the **nn** package




the **nn** package

- nn: neural networks made easy
- building blocks of differentiable modules

➔ define a model with pre-normalization, to work on raw RGB images:

```
01 model = nn.Sequential()  
02  
03 model:add( nn.SpatialConvolution(3,16,5,5) )  
04 model:add( nn.Tanh() )  
05 model:add( nn.SpatialMaxPooling(2,2,2,2) )  
06 model:add( nn.SpatialContrastiveNormalization(16, image.gaussian(3)) )  
07  
08 model:add( nn.SpatialConvolution(16,64,5,5) )  
09 model:add( nn.Tanh() )  
10 model:add( nn.SpatialMaxPooling(2,2,2,2) )  
11 model:add( nn.SpatialContrastiveNormalization(64, image.gaussian(3)) )  
12  
13 model:add( nn.SpatialConvolution(64,256,5,5) )  
14 model:add( nn.Tanh() )  
15 model:add( nn.Reshape(256) )  
16 model:add( nn.Linear(256,10) )  
17 model:add( nn.LogSoftMax() )
```






the **nn** package

- ➡ When training neural nets, autoencoders, linear regression, convolutional networks, and any of these models, we're interested in gradients, and loss functions
- ➡ The **nn** package provides a large set of transfer functions, which all come with three methods:
 - ➡ `upgradeOutput()` -- compute the output given the input
 - ➡ `upgradeGradInput()` -- compute the derivative of the loss wrt input
 - ➡ `accGradParameters()` -- compute the derivative of the loss wrt weights
- ➡ The **nn** package provides a set of common loss functions, which all come with two methods:
 - ➡ `upgradeOutput()` -- compute the output given the input
 - ➡ `upgradeGradInput()` -- compute the derivative of the loss wrt input



the **nn** package

Compose networks
like Lego blocks



the **nn** package

CUDA Backend via the `cunn` package

```
01  -- define model
02  model = nn.Sequential()
03  model:add( nn.Linear(100,1000) )
04  model:add( nn.Tanh() )
05  model:add( nn.Linear(1000,10) )
06  model:add( nn.LogSoftMax() )
07
08  -- re-cast model as a CUDA model
09  model:cuda()
10
11  -- define input as a CUDA Tensor
12  input = torch.CudaTensor(100)
13  -- compute model's output (is a CudaTensor as well)
14  output = model:forward(input)
15
16  -- alternative: convert an existing DoubleTensor to a CudaTensor:
17  input = torch.randn(100):cuda()
18  output = model:forward(input)
```



the **nngraph** package

Graph composition using chaining

```
In [ ]: -- it is common style to mark inputs with identity nodes for clarity.
input = nn.Identity()()

-- each hidden layer is achieved by connecting the previous one
-- here we define a single hidden layer network
h1 = nn.Tanh()(nn.Linear(20, 10)(input))
output = nn.Linear(10, 1)(h1)
mlp = nn.gModule({input}, {output})

x = torch.rand(20)
dx = torch.rand(1)
mlp:updateOutput(x)
mlp:updateGradInput(x, dx)
mlp:accGradParameters(x, dx)

-- draw graph (the forward graph, '.fg')
-- this will produce an SVG in the runtime directory
graph.dot(mlp.fg, 'MLP', 'MLP')
itorch.image('MLP.svg')
```



the **optim** package

A purely functional view of the world

```
config = {  
  learningRate = 1e-3,  
  momentum = 0.5  
}  
  
for i, sample in ipairs(training_samples) do  
  local func = function(x)  
    -- define eval function  
    return f, df_dx  
  end  
  optim.sgd(func, x, config)  
end
```



the **optim** package


- *Stochastic Gradient Descent*
- *Averaged Stochastic Gradient Descent*
- *L-BFGS*
- *Congugate Gradients*
- *AdaDelta*
- *AdaGrad*
- *Adam*
- *AdaMax*
- *FISTA with backtracking line search*
- *Nesterov's Accelerated Gradient method*
- *RMSprop*
- *Rprop*
- *CMAES*



the **optim** package

Collecting the parameters of your neural net

- Substitute each module weights and biases by one large tensor, making weights and biases point to parts of this tensor



the **optim** package

A purely functional view of the world

```
config = {  
  learningRate = 1e-3,  
  momentum = 0.5  
}  
  
for i, sample in ipairs(training_samples) do  
  local func = function(x)  
    -- define eval function  
    return f, df_dx  
  end  
  optim.sgd(func, x, config)  
end
```



the **optim** package

Define closure

```
08 -- define a closure, that computes the loss, and dloss/dx
09 feval = function()
10   -- select a new training sample
11   _nidx_ = (_nidx_ or 0) + 1
12   if _nidx_ > (#data)[1] then _nidx_ = 1 end
13
14   local sample = data[_nidx_]
15   local inputs = sample[1]
16   local target = sample[2]
17
18   -- reset gradients (gradients are always accumulated,
19   --                      to accomodate batch methods)
20   dl_dx:zero()
21
22   -- evaluate the loss function and its derivative wrt x,
23   -- for that sample
24   local loss_x = criterion:forward(model:forward(inputs), target)
25   model:backward(inputs, criterion:backward(model.output, target))
26
27   -- return loss(x) and dloss/dx
28   return loss_x, dl_dx
29 end
30
```



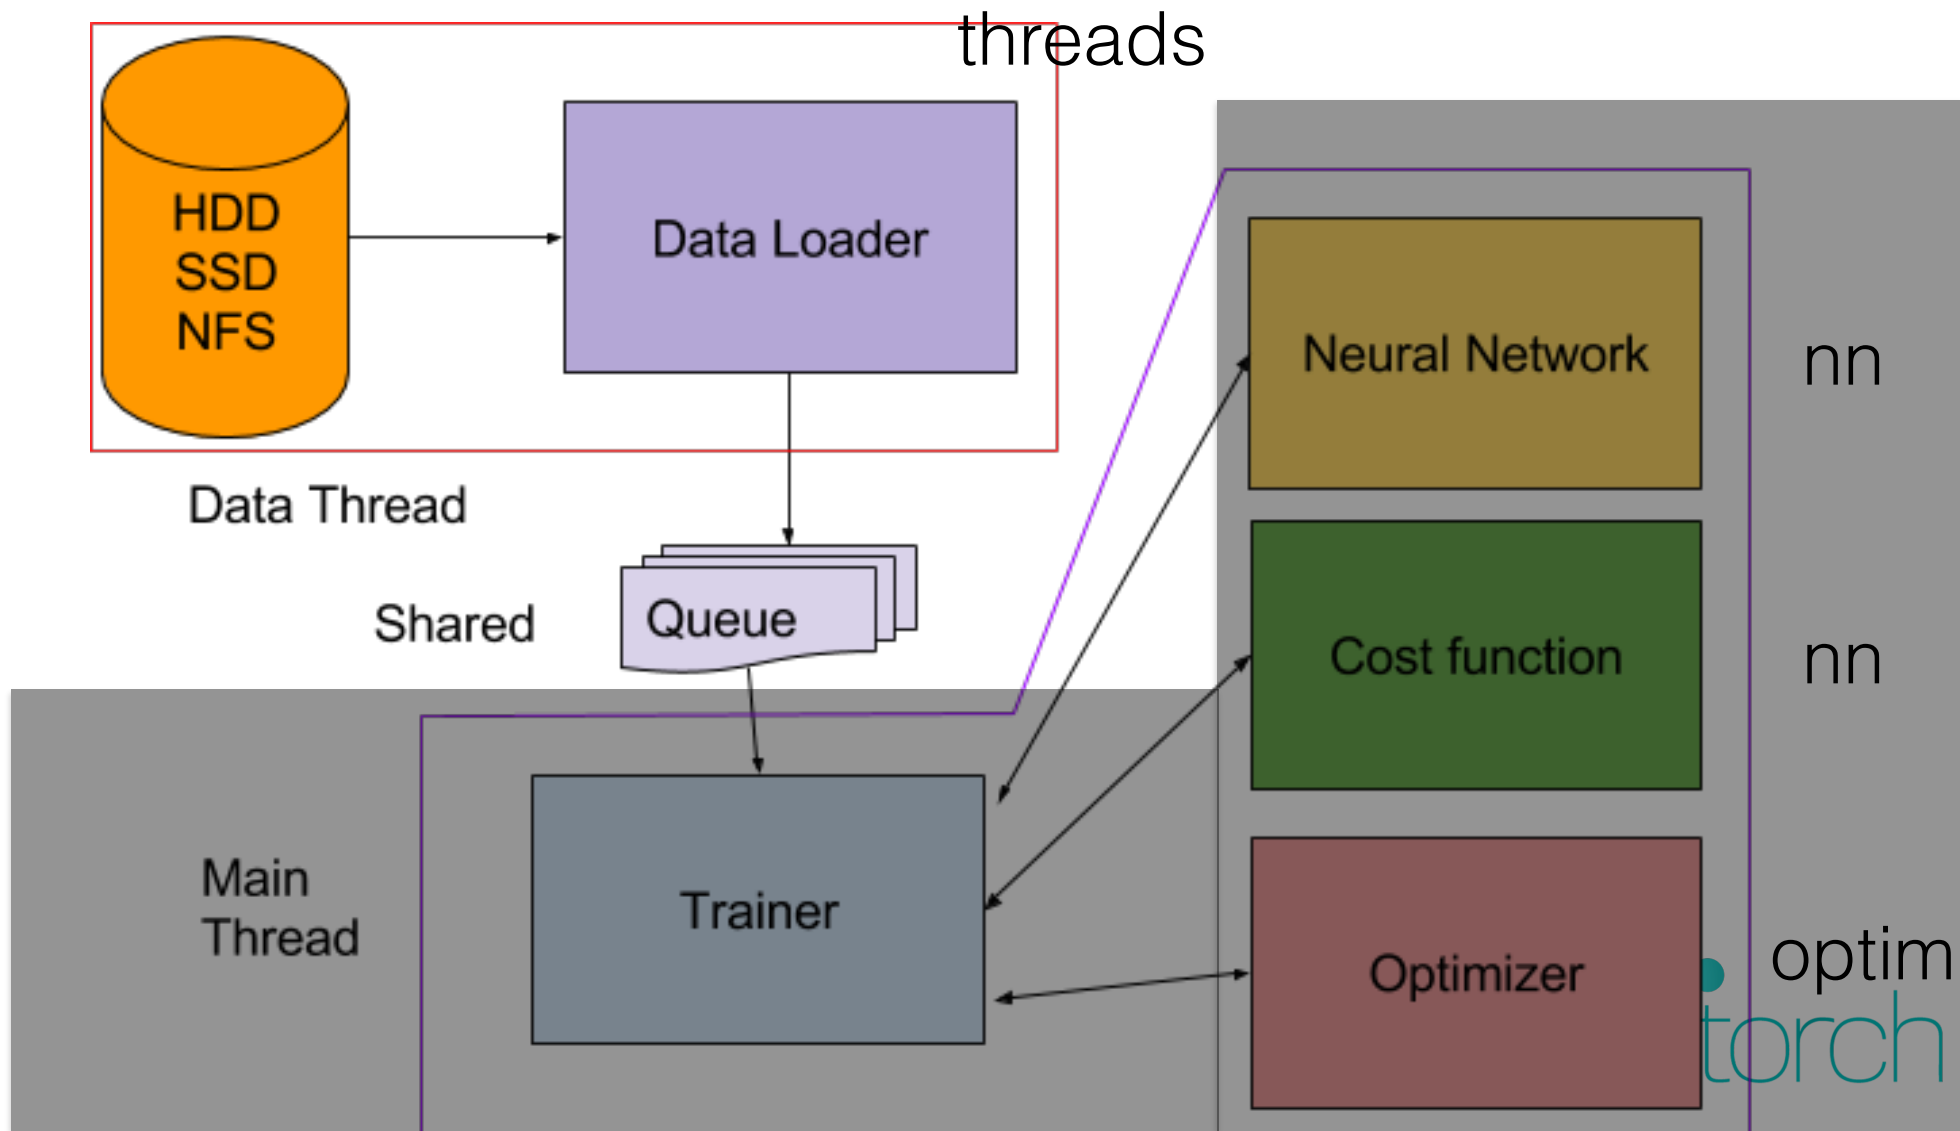
the **optim** package

Define closure

```
31 -- SGD parameters
32 sgd_params = {learningRate = 1e-3, learningRateDecay = 1e-4,
33               weightDecay = 0, momentum = 0}
34
35 -- train for a number of epochs
36 epochs = 1e2
37 for i = 1,epochs do
38   -- this variable is used to estimate the average loss
39   current_loss = 0
40
41   -- an epoch is a full loop over our training data
42   for i = 1,(#data)[1] do
43
44     -- one step of SGD optimization (steepest descent)
45     _,fs = optim.sgd(feval,x,sgd_params)
46
47     -- accumulate error
48     current_loss = current_loss + fs[1]
49   end
50
51   -- report average error on epoch
52   current_loss = current_loss / (#data)[1]
53   print(' current loss = ' .. current_loss)
54 end
```



the **threads** package



the **threads** package

- Create data-loading threads on demand
 - (hilariously called donkeys and it stuck)
- callbacks that are executed in the main thread get data from donkeys, and call optimization functions

```
local ffi = require 'ffi'
local Threads = require 'threads'
Threads.serialization('threads.sharedserialize')

-- This script contains the logic to create K threads for parallel data-loading.
-- For the data-loading details, look at donkey.lua
-----
do -- start K datathreads (donkeys)
  if opt.nDonkeys > 0 then
    local options = opt -- make an upvalue to serialize over to donkey threads
    donkeys = Threads(
      opt.nDonkeys,
      function()
        require 'torch'
      end,
      function(idx)
        opt = options -- pass to all donkeys via upvalue
        tid = idx
        local seed = opt.manualSeed + idx
        torch.manualSeed(seed)
        print(string.format('Starting donkey with id: %d seed: %d', tid, seed))
        paths.dofile('donkey.lua')
      end
    );
```



the **threads** package

```
for i=1,opt.epochSize do
    -- queue jobs to data-workers
    donkeys:addjob(
        -- the job callback (runs in data-worker thread)
        function()
            local inputs, labels = trainLoader:sample(opt.batchSize)
            return inputs, labels
        end,
        -- the end callback (runs in the main thread)
        trainBatch
    )
end

donkeys:synchronize()
```

```
function trainBatch(inputsCPU, labelsCPU)
    cutorch:synchronize()
    collectgarbage()
    local dataLoadingTime = dataTimer:time().real
    timer:reset()

    -- transfer over to GPU
    inputs:resize(inputsCPU:size()):copy(inputsCPU)
    labels:resize(labelsCPU:size()):copy(labelsCPU)

    local err, outputs
    feval = function(x)
        model:zeroGradParameters()
        outputs = model:forward(inputs)
        err = criterion:forward(outputs, labels)
        local gradOutputs = criterion:backward(outputs, labels)
        model:backward(inputs, gradOutputs)
        return err, gradParameters
    end
    optim.sgd(feval, parameters, optimState)
```



Next up

- A complete example of using nn + optim + threads for image generation
- the magic autograd package
- torchnet: common patterns for Torch by Facebook

