



Extensions

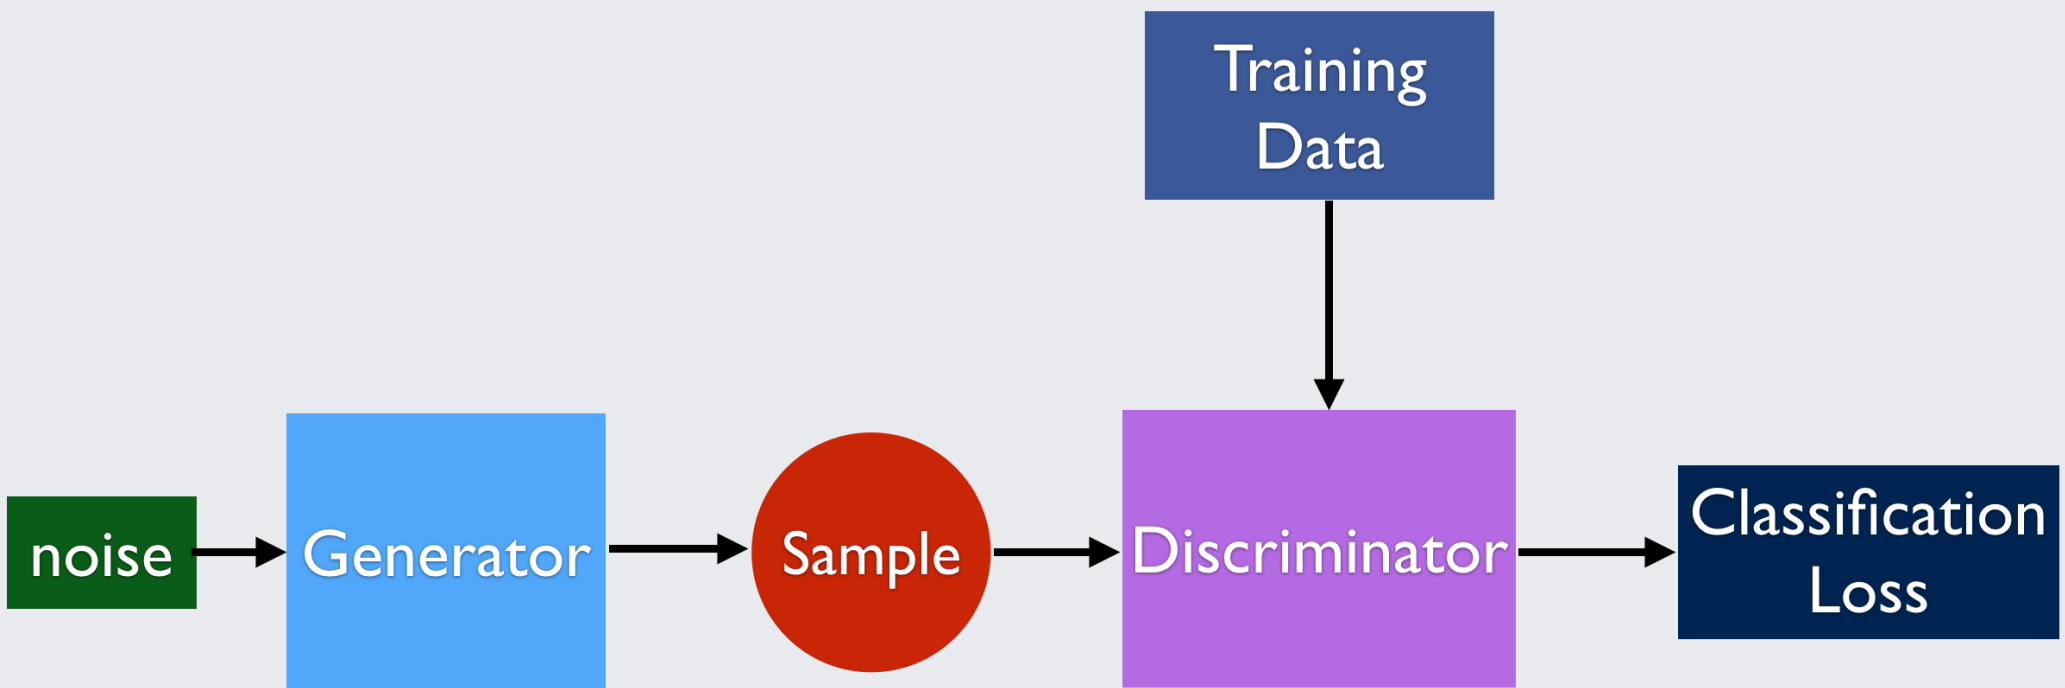
Soumith Chintala
Facebook AI Research

Overview

- A complete example of using nn + optim + threads for image generation
- the magic autograd package
- torchnet: common patterns for Torch by Facebook

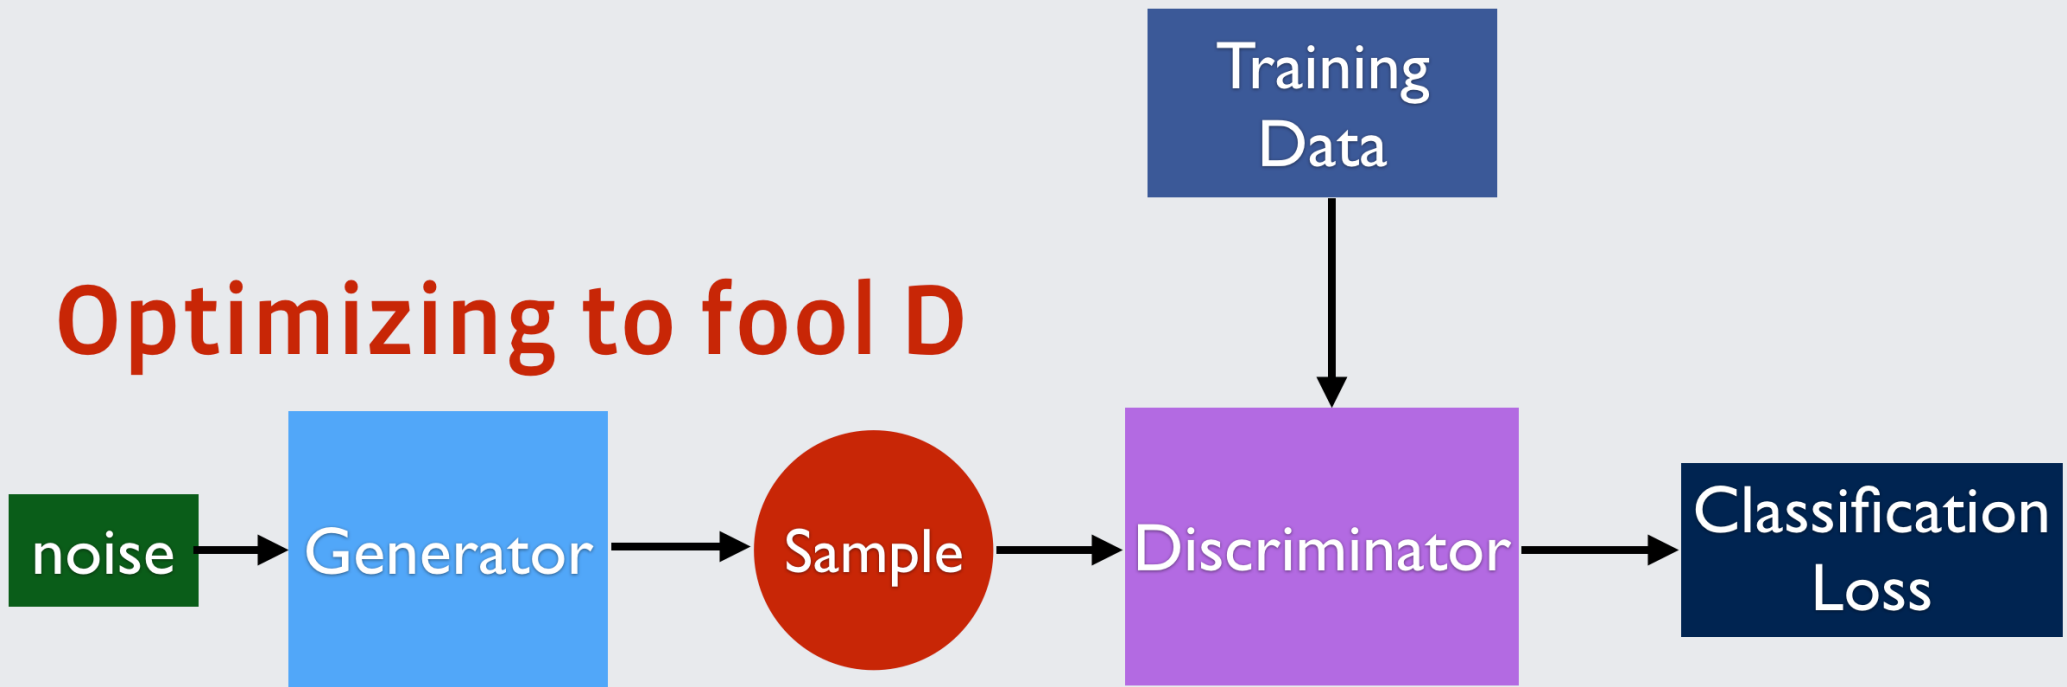


DCGAN

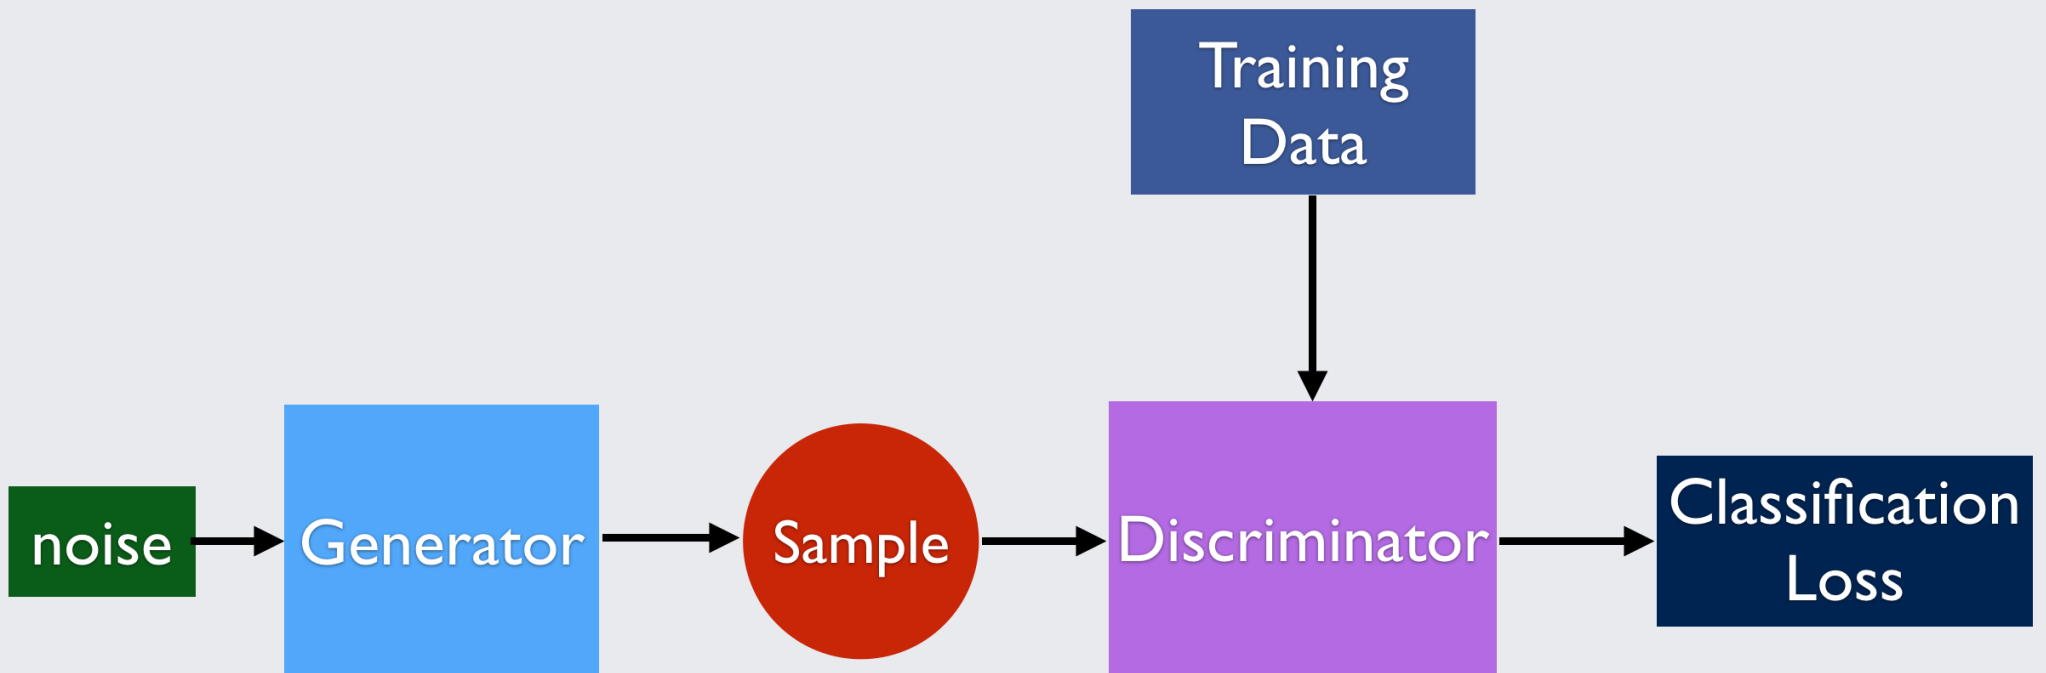


DCGAN

Optimizing to fool D

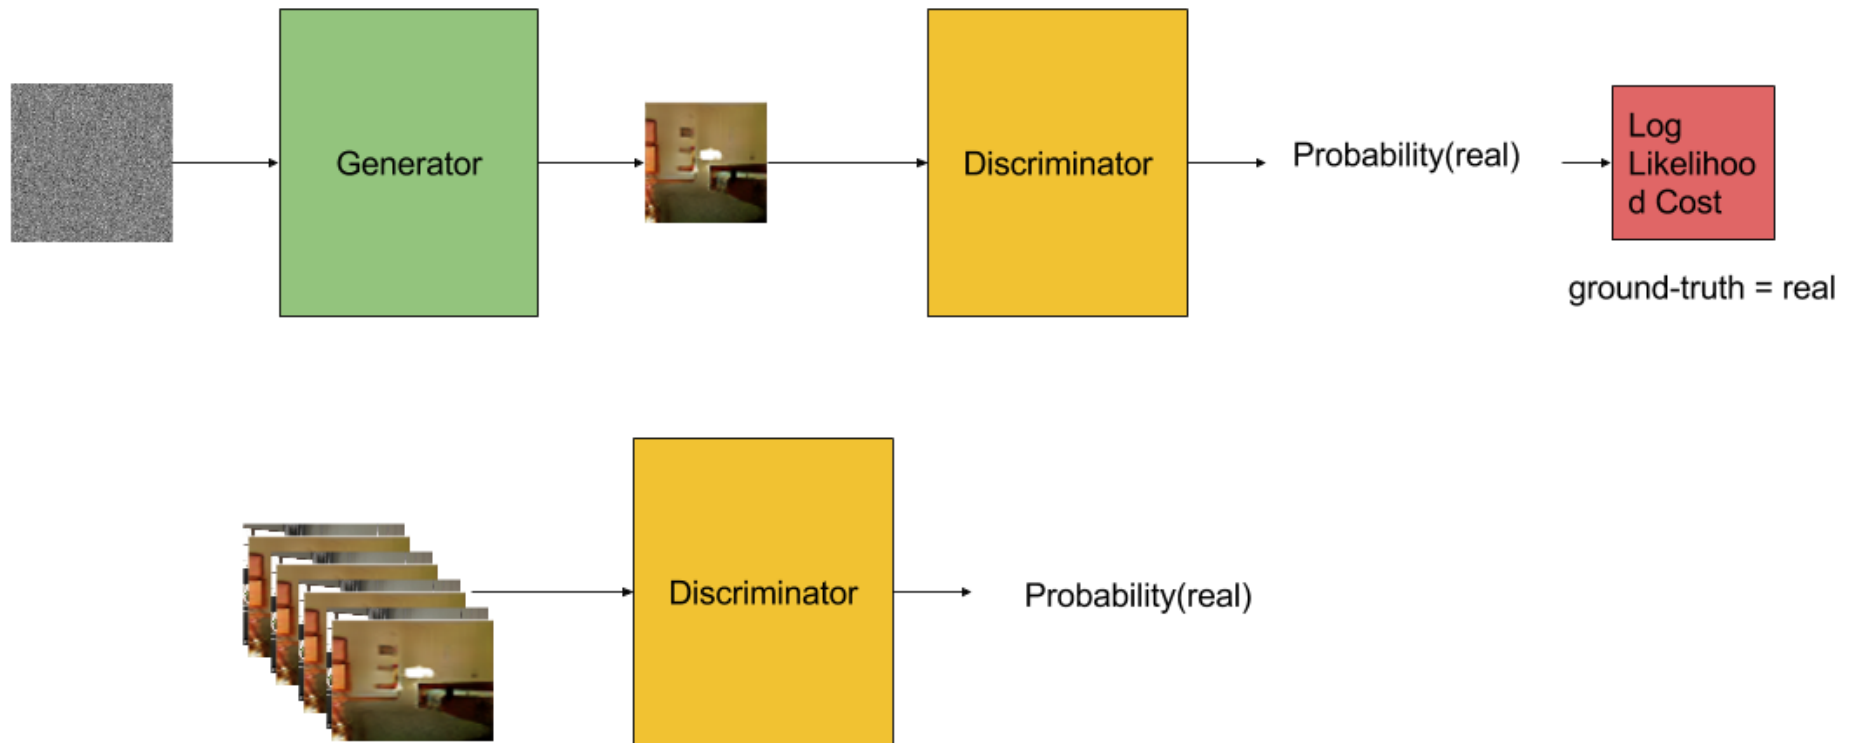


DCGAN



Optimizing to not get fooled by G

DCGAN



DCGAN

Boilerplate

```
1  require 'torch'
2  require 'nn'
3  require 'optim'
4  require 'xlua'
5
6  opt = {
7      batchSize = 64,
8      loadSize = 96,
9      fineSize = 64,
10     nz = 100,           -- # of dim for Z
11     ngf = 64,           -- # of gen filters in first conv layer
12     ndf = 64,           -- # of discrim filters in first conv layer
13     nThreads = 4,       -- # of data loading threads to use
14     niter = 25,          -- # of iter at starting learning rate
15     gpu = 1,             -- gpu = 0 is CPU mode. gpu=X is GPU mode on GPU X
16     name = 'experiment1',
17 }
18
19 xlua.envparams(opt)
20
21 -- create data loader
22 local DataLoader = paths.dofile('data/data.lua')
23 local data = DataLoader.new(opt.nThreads, opt.dataset, opt)
24 print("Dataset: " .. opt.dataset, " Size: ", data:size())
```



DCGAN

Generator

```
33 local netG = nn.Sequential()
34 -- input is Z, going into a convolution
35 netG:add(nn.SpatialFullConvolution(nz, ngf * 8, 4, 4))
36 netG:add(nn.SpatialBatchNormalization(ngf * 8)):add(nn.ReLU(true))
37 -- state size: (ngf*8) x 4 x 4
38 netG:add(nn.SpatialFullConvolution(ngf * 8, ngf * 4, 4, 4, 2, 2, 1, 1))
39 netG:add(nn.SpatialBatchNormalization(ngf * 4)):add(nn.ReLU(true))
40 -- state size: (ngf*4) x 8 x 8
41 netG:add(nn.SpatialFullConvolution(ngf * 4, ngf * 2, 4, 4, 2, 2, 1, 1))
42 netG:add(nn.SpatialBatchNormalization(ngf * 2)):add(nn.ReLU(true))
43 -- state size: (ngf*2) x 16 x 16
44 netG:add(nn.SpatialFullConvolution(ngf * 2, ngf, 4, 4, 2, 2, 1, 1))
45 netG:add(nn.SpatialBatchNormalization(ngf)):add(nn.ReLU(true))
46 -- state size: (ngf) x 32 x 32
47 netG:add(nn.SpatialFullConvolution(ngf, nc, 4, 4, 2, 2, 1, 1))
48 netG:add(nn.Tanh())
49 -- state size: (nc) x 64 x 64
```



DCGAN

Discriminator

```
51 local netD = nn.Sequential()
52
53 -- input is (nc) x 64 x 64
54 netD:add(nn.SpatialConvolution(nc, ndf, 4, 4, 2, 2, 1, 1))
55 netD:add(nn.LeakyReLU(0.2, true))
56 -- state size: (ndf) x 32 x 32
57 netD:add(nn.SpatialConvolution(ndf, ndf * 2, 4, 4, 2, 2, 1, 1))
58 netD:add(nn.SpatialBatchNormalization(ndf * 2)):add(nn.LeakyReLU(0.2, true))
59 -- state size: (ndf*2) x 16 x 16
60 netD:add(nn.SpatialConvolution(ndf * 2, ndf * 4, 4, 4, 2, 2, 1, 1))
61 netD:add(nn.SpatialBatchNormalization(ndf * 4)):add(nn.LeakyReLU(0.2, true))
62 -- state size: (ndf*4) x 8 x 8
63 netD:add(nn.SpatialConvolution(ndf * 4, ndf * 8, 4, 4, 2, 2, 1, 1))
64 netD:add(nn.SpatialBatchNormalization(ndf * 8)):add(nn.LeakyReLU(0.2, true))
65 -- state size: (ndf*8) x 4 x 4
66 netD:add(nn.SpatialConvolution(ndf * 8, 1, 4, 4))
67 netD:add(nn.Sigmoid())
68 -- state size: 1 x 1 x 1
69 netD:add(nn.View(1):setNumInputDims(3))
70 -- state size: 1
```



DCGAN

Loss

```
72 | local criterion = nn.BCECriterion()
```



DCGAN

Optim configuration

```
74  optimStateG = {  
75      learningRate = 0.0002,  
76      beta1 = 0.5,  
77  }  
78  optimStateD = {  
79      learningRate = opt.lr,  
80      beta1 = 0.5,  
81  }
```

DCGAN

Create buffers and push them to GPU


```
82 -----
83 local input = torch.Tensor(opt.batchSize, 3, opt.fineSize, opt.fineSize)
84 local noise = torch.Tensor(opt.batchSize, nz, 1, 1)
85 local label = torch.Tensor(opt.batchSize)
86 local errD, errG
87 -----
88 if opt.gpu > 0 then
89     require 'cunn'
90     cutorch.setDevice(opt.gpu)
91     input = input:cuda(); noise = noise:cuda(); label = label:cuda()
92     netD:cuda(); netG:cuda(); criterion:cuda()
93 end
```



DCGAN

flatten the parameters

```
95 | local parametersD, gradParametersD = netD:getParameters()  
96 | local parametersG, gradParametersG = netG:getParameters()
```



DCGAN

Define $y = D(w_d, \text{input})$

```
98  -- create closure to evaluate f(X) and df/dX of discriminator
99  local fDx = function(x)
100    gradParametersD:zero()
101
102    -- train with real
103    local real = data:getBatch()
104    input:copy(real)
105    label:fill(real_label)
106
107    local output = netD:forward(input)
108    local errD_real = criterion:forward(output, label)
109    local df_do = criterion:backward(output, label)
110    netD:backward(input, df_do)
111
112    -- train with fake
113    noise:uniform(-1, 1)
114    local fake = netG:forward(noise)
115    input:copy(fake)
116    label:fill(fake_label)
117
118    local output = netD:forward(input)
119    local errD_fake = criterion:forward(output, label)
120    local df_do = criterion:backward(output, label)
121    netD:backward(input, df_do)
122
123    errD = errD_real + errD_fake
124
125    return errD, gradParametersD
126  end
```



DCGAN

Define $y = G(w_g, z)$

```
128 -- create closure to evaluate f(X) and df/dX of generator
129 local fGx = function(x)
130     gradParametersG:zero()
131
132     noise:uniform(-1, 1) -- regenerate random noise
133     local fake = netG:forward(noise)
134     input:copy(fake)
135     label:fill(real_label) -- fake labels are real for generator cost
136
137     local output = netD:forward(input)
138     errG = criterion:forward(output, label)
139     local df_do = criterion:backward(output, label)
140     local df_dg = netD:updateGradInput(input, df_do)
141
142     netG:backward(noise, df_dg)
143     return errG, gradParametersG
144 end
```



DCGAN

Now train!

```
146 -- train
147 for epoch = 1, opt.niter do
148     for i = 1, data:size(), opt.batchSize do
149         -- (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
150         optim.adam(fDx, parametersD, optimStateD)
151
152         -- (2) Update G network: maximize log(D(G(z)))
153         optim.adam(fGx, parametersG, optimStateG)
154     end
155     paths.mkdir('checkpoints')
156     torch.save('checkpoints/' .. opt.name .. '_' .. epoch .. '_net_G.t7', netG)
157     torch.save('checkpoints/' .. opt.name .. '_' .. epoch .. '_net_D.t7', netD)
158 end
```



DCGAN

Now train!

[VIDEO REMOVED]

DCGAN

Generations!

```
local images = net:forward(noise)
```

Noise:

DCGAN

Generations!

```
local images = net:forward(noise)
```

Noise:

```
noise = torch.Tensor(opt.batchSize, opt.nz, opt.imsz, opt.imsz)  
noise:uniform(-1, 1)
```



DCGAN

Generations!

```
local images = net:forward(noise)
```

Interpolations in z space:

```
noise = torch.Tensor(opt.batchSize, opt.nz, opt.imsz, opt.imsz)
noiseL = torch.FloatTensor(opt.nz):uniform(-1, 1)
noiseR = torch.FloatTensor(opt.nz):uniform(-1, 1)
-- do a linear interpolation in Z space between point A and point B
-- each sample in the mini-batch is a point on the line
line = torch.linspace(0, 1, opt.batchSize)
for i = 1, opt.batchSize do
    noise:select(1, i):copy(noiseL * line[i] + noiseR * (1 - line[i]))
end
```




DCGAN

Generations!

Interpolations

```
noise = torch.randn(1, 1, 1, 1)  
noiseL = torch.randn(1, 1, 1, 1)  
noiseR = torch.randn(1, 1, 1, 1)  
-- do a linear interpolation  
-- each sample is a line  
line = torch.randn(1, 1, 1, 1)  
for i = 1, opt.n_samples  
  noise:select(1, i)  
end
```



```
e, opt.imsz)
```

and point B

```
(1 - line[i]))
```



DCGAN

Generations!

```
local images = net:forward(noise)
```

Arithmetic in z space:

<https://github.com/soumith/dcgan.torch/blob/simple/arithmetic.lua>

LIVE DEMO



autograd by

Basic usage

```
In [1]: autograd = require 'autograd'
```

```
In [38]: foo = function(a, b, c)
          return a + b * c
        end
```

```
In [39]: dfoo = autograd(foo)
```

```
In [40]: da, f_abc = dfoo(3.5, 2.1, 1.1)
          print('Value: ' .. f_abc, 'Gradient: ', da)
```

```
Out[40]: Value: 5.81      Gradient:      1
```



autograd by

Conditionals

```
In [41]: foo = function(a, b, c)
          if b > c then
              return a * math.sin(b)
          else
              return a + b * c
          end
        end
```

```
In [42]: dfoo = autograd(foo)
```

```
In [43]: da, f_abc = dfoo(3.5, 2.1, 1.1)
          print('Value: ' .. f_abc, 'Gradient: ', da)
```

```
Out[43]: Value: 3.0212327832711 Gradient: 0.86320936664887
```

```
In [44]: da, f_abc = dfoo(3.5, 0.1, 1.1)
          print('Value: ' .. f_abc, 'Gradient: ', da)
```

```
Out[44]: Value: 3.61 Gradient: 1
```



autograd by

Dynamic “while” evaluations

```
In [45]: foo = function(a, b, c)
          local val = 0
          while b > c do
            val = val + a * math.sin(b)
            b = b - 0.1
          end
          return val
        end
```

```
In [46]: dfoo = autograd(foo)
```

```
In [47]: da, f_abc = dfoo(3.5, 2.1, 1.1)
          print('Value: ' .. f_abc, 'Gradient: ', da)
```

```
Out[47]: Value: 33.468522197289   Gradient:      9.5624349135111
```



autograd by

Dependency checks

```
In [45]: foo = function(a, b, c)
          local val = 0
          while b > c do
            val = val + a * math.sin(b)
            b = b - 0.1
          end
          return val
        end
```

```
In [46]: dfoo = autograd(foo)
```

```
In [47]: da, f_abc = dfoo(3.5, 2.1, 1.1)
          print('Value: ' .. f_abc, 'Gradient: ', da)
```

```
Out[47]: Value: 33.468522197289 Gradient: 9.5624349135111
```

```
[48]: da, f_abc = dfoo(3.5, 1.0, 1.1)
       print('Value: ' .. f_abc, 'Gradient: ', da)
```

...all/share/luarocks-5.1/autograd/runtime/direct/DirectTape.lua:119: A node type was not returned. This is either because a gradient was not defined, or the input is independent of the output

autograd by

- Dynamic graphs
- Auto-differentiated nn modules
- Tape based autograd

Torchnet

<https://github.com/torchnet/torchnet>



Next steps

- <http://torch.ch/>
- <https://github.com/szagoruyko/idiap-tutorials>

