

Training ConvNets with Torch

Soumith Chintala

Next.ml

17th January, 2014

Tutorial Structure

Hour 1

- Action Recognition from Videos:
 - Discuss task at hand
 - End-to-end rule-based methods
 - Hand-engineered features + SVM
 - Neural Nets (and discuss discovering graph structure automatically)
 - ConvNets
- Notebook Setup

Tutorial Structure

Hour 2

- Torch Basics
 - What is torch?
 - Lua and Syntax
 - Tensors, types, storages
 - Slicing and selecting
 - views vs copies
- Basic visualization in iTorch
- Neural Networks
 - Containers and Modules
 - Examples:
 - Overfeat
 - Alexnet
 - Inception
 - LSTM
- Optim package

Tutorial Structure

- Hour 2/3
- Coding an action recognition system
 - load and normalize data
 - create neural network
 - train neural network
 - test/validate neural network

Task at Hand

First attempt

Pure rule-based

Second Attempt

- Hand-designed features

Third Attempt

- Convolution Networks


Fourth Attempt

- Fully connected nets

Neural net in detail

- Show a typical two-layer network

●●



torch

What is Torch?

- Not a language. Based on Lua
- Not just a Tensor library
- Ecosystem and community

Lua and Syntax

Tensors

- nDimensional arrays – `torch.Tensor`
- Separated from Storages
- Can be sliced, re-viewed, selected etc. without a memory copy where possible

Basic Visualization

Neural Networks

Module

A neural network is called a **Module** (or simply *module* in this documentation) in Torch. `Module` is an abstract class which defines four main methods:

- `forward(input)` which computes the output of the module given the `input` `Tensor`.
- `backward(input, gradOutput)` which computes the gradients of the module with respect to its own parameters, and its own inputs.
- `zeroGradParameters()` which zeroes the gradient with respect to the parameters of the module.
- `updateParameters(learningRate)` which updates the parameters after one has computed the gradients with `backward()`

It also declares two members:

- `output` which is the output returned by `forward()` .
- `gradInput` which contains the gradients with respect to the input of the module, computed in a `backward()` .

Two other perhaps less used but handy methods are also defined:

- `share(mlp,s1,s2,...,sn)` which makes this module share the parameters `s1,..sn` of the module `mlp` . This is useful if you want to have modules that share the same weights.
- `clone(...)` which produces a deep copy of (i.e. not just a pointer to) this Module, including the current state of its parameters (if any).

Neural Networks

Plug and play

Building a simple neural network can be achieved by constructing an available layer. A linear neural network (perceptron!) is built only in one line:

```
mlp = nn.Linear(10,1) -- perceptron with 10 inputs
```

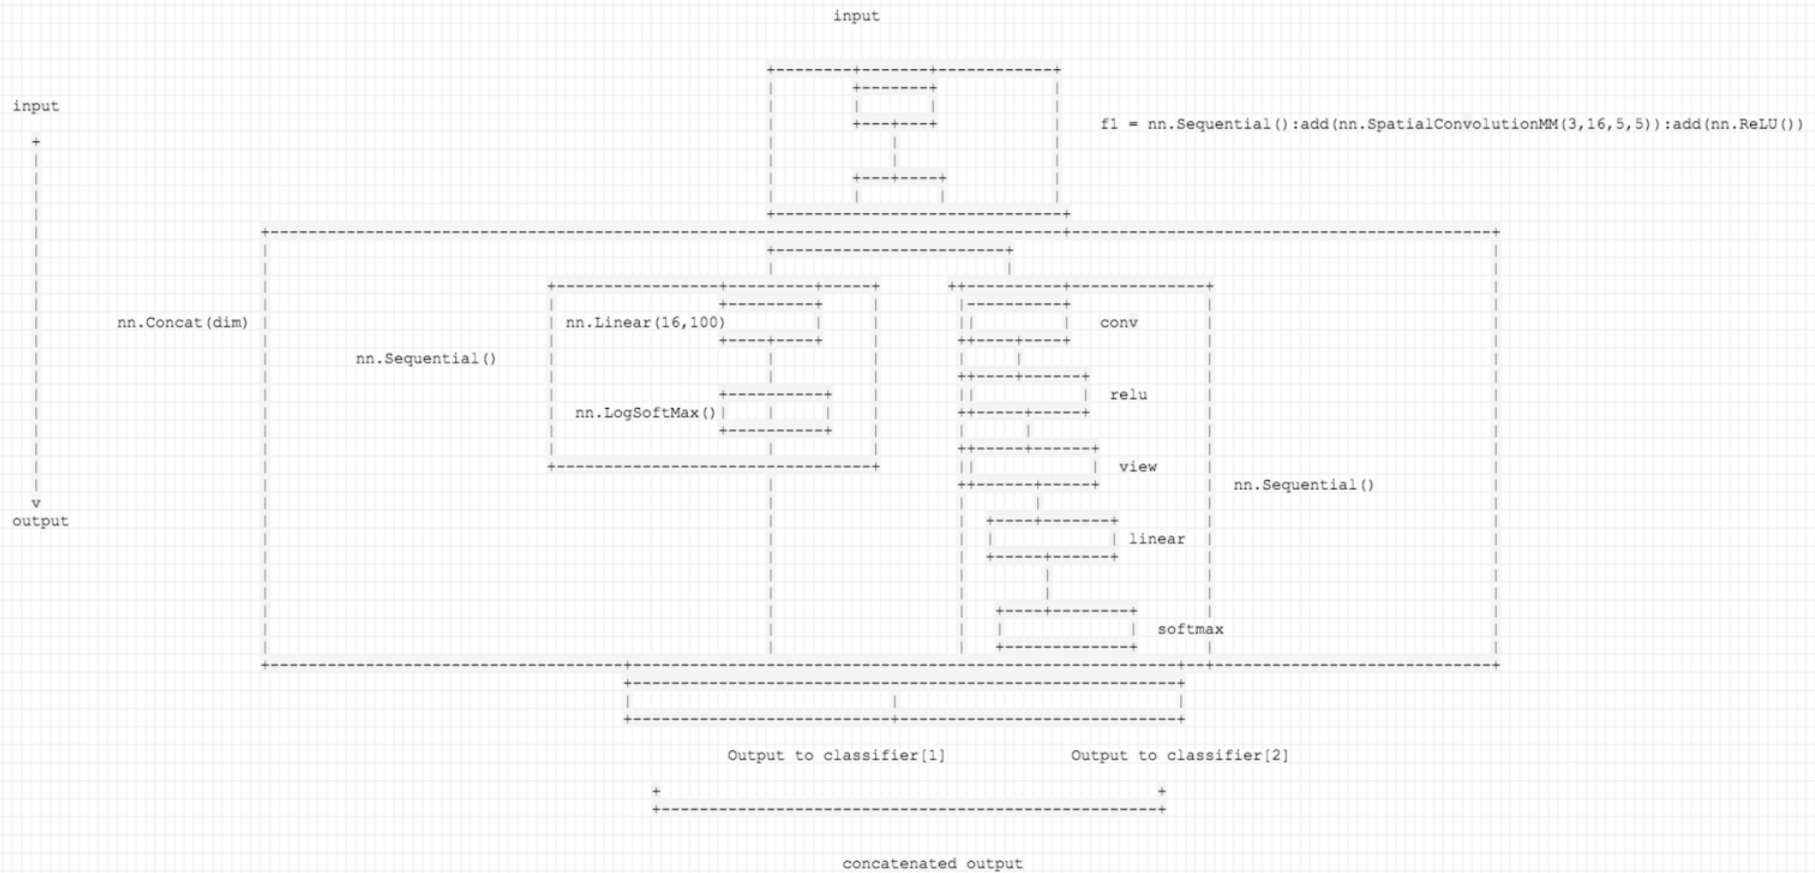
More complex neural networks are easily built using container classes `Sequential` and `Concat`. `Sequential` plugs layer in a feed-forward fully connected manner. `Concat` concatenates in one layer several modules: they take the same inputs, and their output is concatenated.

Creating a one hidden-layer multi-layer perceptron is thus just as easy as:

```
mlp = nn.Sequential()  
mlp.add( nn.Linear(10, 25) ) -- 10 input, 25 hidden units  
mlp.add( nn.Tanh() ) -- some hyperbolic tangent transfer function  
mlp.add( nn.Linear(25, 1) ) -- 1 output
```

Of course, `Sequential` and `Concat` can contains other `Sequential` or `Concat` , allowing you to try the craziest neural networks you ever dreamt of! See the `[[#nn.Modules|complete list of available modules]]`.

Neural Networks



Optimization

We first define a state for conjugate gradient:

```
state = {  
    verbose = true,  
    maxIter = 100  
}
```

and now we train:

```
x = torch.rand(N)  
optim.cg(JdJ, x, state)
```

You should see something like:

```
after 120 evaluation J(x) = -3.136835  
after 121 evaluation J(x) = -3.136836  
after 122 evaluation J(x) = -3.136837  
after 123 evaluation J(x) = -3.136838
```